

Certification of Imitation Controllers Using Learned Lyapunov Functions

Master thesis by Josua Christoph Lindemann

Date: August 18, 2026, Date of submission: 16. August 2026

1. Review: Prof. Dr.-Ing. Rolf Findeisen

2. Review: Hendrik Alsmeier M.Sc.

Darmstadt



TECHNISCHE
UNIVERSITÄT
DARMSTADT

CCPS

Control and Cyber-Physical Systems

Electrical Engineering and
Information Technology
Department

Certification of Imitation Controllers Using Learned Lyapunov Functions

Master thesis by Josua Christoph Lindemann

1. Review: Prof. Dr.-Ing. Rolf Findeisen
2. Review: Hendrik Alsmeier M.Sc.

Date: August 18, 2026

Date of submission: 16. August 2026

Darmstadt

Technische Universität Darmstadt
Institut für Automatisierungstechnik und Mechatronik
Fachgebiet Control and Cyber-Physical Systems
Prof. Dr.-Ing. Rolf Findeisen

Contents

1	Abstract	1
2	Introduction	2
3	Optimal Control and Lyapunov Stability Theory	3
3.1	Discrete-Time Dynamical Systems and Stability	3
3.1.1	Lyapunov Stability	3
3.2	Model Predictive Control and Stability	4
3.2.1	Problem Formulation	4
3.2.2	Terminal Ingredients and their Construction	5
3.2.3	Recursive Feasibility and Closed-Loop Stability	6
4	Foundations of Neural Control Policies	9
4.1	Artificial Neural Networks	9
4.1.1	Multi-Layer Perceptrons	9
4.1.2	Optimization and Training	11
4.2	Adversarial Examples and Projected Gradient Methods	13
4.2.1	Projected Gradient Descent	13
4.2.2	Adaptation to Neural Lyapunov Control	14
5	Formal Verification of Neural Control Policies	15
5.1	Incomplete Verification	15
5.1.1	CROWN	15
5.1.2	α -CROWN	15
5.2	Complete Verification	15
5.2.1	β -CROWN	15
6	Learning Lyapunov-Stable Imitation Controllers	16
6.1	Neural Policy Approximation and Imitation Learning	16
6.2	Lyapunov Learning for Stability Certification	18
6.2.1	Lyapunov Network Architecture	18
6.2.2	Loss Formulation	19
6.2.3	Falsification-Guided Training	23
6.2.4	ρ -Sublevel Estimation	25
6.2.5	Overall Training Procedure	25
6.3	Formal Verification Methodology	27

7	Experimental Evaluation and Certification Results	30
7.1	System Modeling and Problem Formulation	30
7.1.1	Double Integrator	30
7.1.2	Cartpole	30
7.2	Linear System Results: Double Integrator	31
7.3	Nonlinear System Results: Cartpole	31
8	Discussion	33
8.1	Qualitative Discussion of Results	33
8.2	Challenges with Nonlinear Dynamics	33
8.3	Scalability and Limitations	33
9	Conclusion	34
A	Appendix	35
A.0.1	Imitation Learning	35
	Bibliography	37

List of Variables and Symbols

Symbol	Description
<i>System Dynamics & Control</i>	
n_x	State space dimension
n_u	Control input space dimension
$\mathcal{X}$	Admissible state set
$\mathcal{U}$	Admissible input set
$\underline{x}$	Lower state bound
$\bar{x}$	Upper state bound
$\underline{u}$	Lower control input bound
$\bar{u}$	Upper control input bound
T_s	Sampling time
k	Discrete time step index
$x(t)$	Continuous-time state vector
x_k	Discrete-time state vector
$u(t)$	Continuous-time control input vector
u_k	Discrete-time control input vector
x^*	System equilibrium point
u^*	Input reference at x^*
x_{k+1}	Discrete-time successor state
x^+	One-step successor state
$\pi(x)$	State-feedback policy
$f_c(x, u)$	Continuous-time system dynamics
$f(x_k, u_k)$	Discrete-time system dynamics
<i>Optimal Control & Model Predictive Control (MPC)</i>	
N	Prediction horizon
$\ell(x, u)$	Stage cost function
$V_f(x)$	Terminal cost function
$\mathcal{X}_f$	Terminal constraint set
Q	State weighting matrix
R	Input weighting matrix
u_0^*	Optimal control input
$V_N^*(x)$	Optimal value function
A	Linearized system state matrix
B	Linearized system input matrix
K	LQR feedback gain matrix

Symbol	Description
P	DARE solution matrix
α_N	Relaxed Lyapunov decrease factor
$\mathcal{X}_N$	Feasible initial state set
<i>Machine Learning</i>	
L	Total layer count
l	Layer index
n_l	Layer l neuron count
n_0	Input layer dimension
n_L	Output layer dimension
$a^{(l)}$	Layer activation vector
$z^{(l)}$	Layer pre-activation vector
$W^{(l)}$	Layer weight matrix
$b^{(l)}$	Layer bias vector
$\sigma^{(l)}(\cdot)$	Non-linear activation function
$\tanh(\cdot)$	Hyperbolic tangent activation function
$[z]_+$	Rectified linear unit (ReLU) activation
$\hat{y}$	Network prediction vector
y	Target output vector
θ	Trainable network parameters
θ^*	Optimal network parameters
$\mathcal{D}$	Training dataset
M	Training data size
$\mathcal{B}$	Mini-batch
B	Mini-batch size
$\mathcal{L}$	Total training loss function
$\mathcal{L}_{\text{MSE}}$	Mean squared error loss
$\mathcal{L}_{\text{reg}}$	Regularization loss
$\mathcal{L}_{\text{aux}}$	Auxiliary loss
ω_{reg}	Regularization weight
ω_{aux}	Auxiliary weight
$\delta^{(l)}$	Loss sensitivity per layer
$\delta^{(L)}$	Loss sensitivity at last layer
$\odot$	Element-wise (Hadamard) product
g	Mini-batch loss gradient
α	Learning rate
E	Total epoch count
t	Optimization timestep index
m_t	Biased first moment estimate
v_t	Biased second moment estimate
$\hat{m}_t$	Bias-corrected first moment estimate
$\hat{v}_t$	Bias-corrected second moment estimate
β_1, β_2	Moment exponential decay rates
ϵ	Optimizer numerical stability constant
$x_{\text{adv}}^{(j)}$	Adversarial candidate state at PGD step j

Symbol	Description
j	PGD iteration index
I	Maximum PGD iteration count
α_{PGD}	PGD step size
δ	Adversarial perturbation vector
ε_δ	Maximum perturbation magnitude
p	Norm order
$\mathcal{P}$	Admissible perturbation set
$\mathcal{J}(x, y; \theta)$	Adversarial loss function
$\Pi_{x+\mathcal{P}}(\cdot)$	Projection operator onto perturbation set
$\text{sign}(\cdot)$	Signum function
<i>Imitation Learning</i>	
$\mathcal{X}_\pi$	Imitation learning state space
$\underline{x}_\pi$	Admissible state lower bound
$\bar{x}_\pi$	Admissible state upper bound
$\mathcal{D}_\pi$	Imitation learning dataset
$N_{\mathcal{D}}$	Imitation learning dataset sample count
$\mathcal{B}_\pi$	Imitation learning training batch
M_π	Imitation learning batch size
$\pi_0(x)$	Expert MPC policy
$\pi_\theta(x)$	Learned imitation policy
$h_\theta(x)$	Unconstrained neural network output
$\text{clamp}(\cdot)$	Control input clamping function
S_x	State scaling matrix
S_u	Input scaling matrix
x_{scale}	State scaling vector
u_{scale}	Input scaling vector
$d(x)$	Scaled equilibrium distance
$w_x(x)$	Behavioral cloning loss weight
$w_{\min}$	Baseline loss weight
η	Decay rate for state-dependent weight
$\mathcal{L}_{\text{BC}}$	Behavioral cloning loss
$\mathcal{L}_f$	Dynamics-aware constraint loss
ω_{BC}	Behavioral cloning loss weight
ω_f	Dynamics-aware constraint loss weight
$\mathcal{L}_\pi$	Total imitation policy loss
<i>Neural Lyapunov Learning & Falsification</i>	
$\pi_V(x)$	Co-trained imitation policy
$V_\phi(x)$	Neural Lyapunov candidate
ϕ	Neural Lyapunov candidate parameters
ε	Regularization constant
R_ϕ	Learnable transformation matrix
ρ	Target Lyapunov sublevel value
$\rho_{\max}$	Theoretical maximum sublevel bound

Symbol	Description
$\mathcal{V}_\rho$	Certified sublevel set
$\mathcal{X}_V$	Lyapunov training domain
$\partial\mathcal{X}_V$	Boundary of Lyapunov training domain
$\rho_{\min}$	Minimum sublevel target bound
e_V	Lyapunov decrease violation
κ	Target decrease rate
ε_{rel}	Numerical stabilization constant
$\mathcal{L}_{\text{dec}}$	Relative Lyapunov decrease loss
S_{range}	Domain range scaling matrix
$\mathcal{L}_{\text{inv}}$	Relative state-invariance loss
$\mathcal{L}_{\text{cond}}$	Combined condition violation loss
w_ρ	Soft sublevel gating weight
s	Gate sharpness parameter
$\mathcal{L}_{\rho\text{-gated}}$	Soft-gated condition loss
$\mathcal{B}_\partial$	Boundary candidate sample set
$\Pi_{\partial\mathcal{X}_V}$	Boundary projection operator
c_j	Active boundary limit along dimension j
$\mathcal{B}_{\text{ROA}}$	ROA shell sample set
ξ	Inner boundary shell factor
$\mathcal{L}_{\text{ROA}}$	Region of Attraction expansion loss
$e_{V,\text{lower}}$	Signed Lyapunov decrease lower bound
$\mathcal{M}_L$	Region LIRPA condition violation
w_L	LIRPA region weight
α_L	LIRPA weight decay rate
N_L	LIRPA region count
$\mathcal{L}_L$	LIRPA loss
$\mathcal{B}_V$	Lyapunov scaling Sobol sample batch
ℓ_V	Mean Lyapunov log-value
$\ell_{V,0}$	Initial mean Lyapunov log-value
$\mathcal{L}_V$	Lyapunov magnitude scaling loss
$\mathcal{B}_{\text{pol}}$	Policy scaling Sobol sample batch
$\mathcal{L}_{\text{pol}}$	Policy scaling loss
$\ell_{R,0}$	Initial matrix R Frobenius norm
$\mathcal{L}_R$	Matrix R regularization loss
$\mathcal{L}_{\text{total}}$	Total combined Lyapunov loss
$\omega_{\rho\text{-gated}}$	Soft-gated condition loss weight
ω_{ROA}	Region of Attraction expansion loss weight
ω_L	LIRPA loss weight
ω_V	Lyapunov magnitude scaling loss weight
ω_{pol}	Policy scaling loss weight
ω_R	Matrix R regularization loss weight
$\mathcal{C}$	Counterexample pool
$\mathcal{S}$	Explorative pool
$\mathcal{B}_\mathcal{S}$	Explorative pool candidate sample set
$N_\mathcal{S}$	Explorative pool target size

Symbol	Description
x_{adv}	Counterexample state
τ	Counterexample age
τ_{max}	Maximum lifespan
m	Sublevel margin
γ	Overestimation factor
$\tilde{\rho}$	Estimated ROA boundary minimum
v	EMA decay rate
N_{outer}	Outer epoch count
K_{steps}	Inner mini-batch optimization steps per epoch
N_{cex}	Counterexample mining step interval
<i>Formal Verification & Certification</i>	
$\mathcal{B}_C$	Formal verification domain
$\Phi(x; \rho)$	Verification specification formula
Φ_{sub}	Sublevel set exclusion predicate
Φ_{pos}	Positive definiteness predicate
Φ_{inv}	Forward invariance predicate
Φ_{dec}	Strict Lyapunov decrease predicate
$\mathcal{B}_{\text{hole}}$	Equilibrium exclusion region
ρ^*	Maximal verified sublevel value



1 Abstract



2 Introduction

3 Optimal Control and Lyapunov Stability Theory

This chapter provides the theoretical foundation for discrete-time dynamical systems and Model Predictive Control (MPC), which serves as the expert policy for the imitation learning framework used in this work. First, Section 3.1 introduces discrete-time dynamical systems, state and input constraints, and the core principles of Lyapunov stability. Afterwards, Section 3.2 outlines the constrained finite-horizon optimal control problem and details the terminal ingredients required to guarantee asymptotic stability of the closed-loop system.

3.1 Discrete-Time Dynamical Systems and Stability

Implementing optimal control schemes on digital hardware requires a discrete-time representation. Consequently, the continuous-time dynamics

$$\dot{x}(t) = f_c(x(t), u(t)) \quad (3.1)$$

must be discretized, where $x(t) \in \mathbb{R}^{n_x}$ and $u(t) \in \mathbb{R}^{n_u}$ denote the continuous-time state and input vectors, respectively, with state space dimension $n_x \in \mathbb{N}$ and input dimension $n_u \in \mathbb{N}$. In this work, the corresponding discrete-time system

$$x_{k+1} = f(x_k, u_k) \quad (3.2)$$

is obtained by approximating the continuous dynamics using a numerical integration scheme with a constant sampling time $T_s > 0$. Here, $k \in \mathbb{N}_0$ denotes the discrete time step index, such that $x_k := x(kT_s) \in \mathbb{R}^{n_x}$ and $u_k := u(kT_s) \in \mathbb{R}^{n_u}$ represent the state and control input at time step k . The discrete-time function $f : \mathbb{R}^{n_x} \times \mathbb{R}^{n_u} \rightarrow \mathbb{R}^{n_x}$ is evaluated using a numerical integration scheme, such as the forward Euler or a fourth-order Runge-Kutta (RK4) method [1].

To account for physical and operational limitations, the system states and control inputs are restricted to constraint sets $\mathcal{X} := [\underline{x}, \bar{x}] \subseteq \mathbb{R}^{n_x}$ and $\mathcal{U} := [\underline{u}, \bar{u}] \subseteq \mathbb{R}^{n_u}$. Furthermore, without loss of generality, any equilibrium (x^*, u^*) can be shifted to the origin via deviation variables $\tilde{x} := x - x^*$ and $\tilde{u} := u - u^*$. For readability, the tilde notation is subsequently dropped, so that x and u represent the shifted variables, assuming $(x^*, u^*) = (\mathbf{0}, \mathbf{0})$ throughout this work.

3.1.1 Lyapunov Stability

Consider the discrete-time closed-loop system under a state-feedback control policy $\pi : \mathcal{X} \rightarrow \mathcal{U}$. The origin serves as an equilibrium point $x^* = \mathbf{0}$ satisfying $\pi(\mathbf{0}) = \mathbf{0}$ and $f(\mathbf{0}, \mathbf{0}) = \mathbf{0}$. Assuming the state

constraint set $\mathcal{X}$ is positively forward invariant under π , the origin is locally asymptotically stable in $\mathcal{X}$ if there exists a positive definite function $V : \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$ satisfying [2, Theorem 2.13]:

$$\begin{aligned} V(\mathbf{0}) &= 0, \\ V(x) &> 0 \quad \forall x \in \mathcal{X} \setminus \{\mathbf{0}\}, \\ V(f(x, \pi(x))) - V(x) &< 0 \quad \forall x \in \mathcal{X} \setminus \{\mathbf{0}\}. \end{aligned} \tag{3.3}$$

Here, the decrease condition ensures that any trajectory starting in $\mathcal{X}$ remains bounded and decays toward the equilibrium x^* .

look up if all assumptions are named

3.2 Model Predictive Control and Stability

This section summarizes the model predictive control (MPC) framework as well as the conditions that are required to ensure the stability of the control loop in optimal control configurations with a finite time horizon.

3.2.1 Problem Formulation

For the nominal MPC setup considered in this work, the following finite-horizon optimal control problem is solved at each time step k given the current system state x_k [2, Sec. 2.2], [3, Sec. 3.3]:

$$\begin{aligned} \min_{(x_{i|k})_{i=0}^N, (u_{i|k})_{i=0}^{N-1}} \quad & J(x_k, (u_{i|k})_{i=0}^{N-1}) = \sum_{i=0}^{N-1} \ell(x_{i|k}, u_{i|k}) + V_f(x_{N|k}) \\ \text{s.t.} \quad & x_{i+1|k} = f(x_{i|k}, u_{i|k}), \quad i = 0, \dots, N-1, \\ & x_{i|k} \in \mathcal{X}, \quad i = 0, \dots, N-1, \\ & u_{i|k} \in \mathcal{U}, \quad i = 0, \dots, N-1, \\ & x_{N|k} \in \mathcal{X}_f, \\ & x_{0|k} = x_k. \end{aligned} \tag{3.4}$$

Here, $N \in \mathbb{N}$ denotes the prediction horizon, $\ell : \mathcal{X} \times \mathcal{U} \rightarrow \mathbb{R}_{\geq 0}$ is the stage cost function, and $V_f : \mathcal{X}_f \rightarrow \mathbb{R}_{\geq 0}$ represents the terminal cost function, which is defined over the terminal constraint set $\mathcal{X}_f \subseteq \mathcal{X}$. To distinguish optimization from closed-loop variables, $x_{i|k}$ and $u_{i|k}$ denote the predicted state and input at horizon step i , initialized at sampling time k . The sequences $(x_{i|k})_{i=0}^N$ and $(u_{i|k})_{i=0}^{N-1}$ serve as optimization variables constrained by the sets $\mathcal{X}$ and $\mathcal{U}$ [3, Sec. 3.2]. Furthermore, the optimal value function is defined as $V_N(x_k) := J(x_k, (u_{i|k}^*)_{i=0}^{N-1})$.

For this work, the stage cost is chosen as the quadratic cost function with positive definite weighting matrices $Q \succ 0$ and $R \succ 0$, which is defined as follows [2, Sec. 1.3]:

$$\ell(x_{i|k}, u_{i|k}) = x_{i|k}^\top Q x_{i|k} + u_{i|k}^\top R u_{i|k}. \tag{3.5}$$

Following the receding horizon principle, only the first optimal control input $u_{0|k}^*$ of the sequence is applied to the system. This implicitly defines the closed-loop MPC feedback policy $\pi_0 : \mathcal{X}_N \rightarrow \mathcal{U}$ as

$$\pi_0(x_k) := u_{0|k}^*, \quad (3.6)$$

yielding the nominal closed-loop dynamics $x_{k+1} = f(x_k, \pi_0(x_k))$. The optimal control problem (3.4) is then repeated at the subsequent step with the updated state [2, Sec. 2.2], [3, Sec. 3.1].

3.2.2 Terminal Ingredients and their Construction

To establish recursive feasibility and closed-loop stability of nominal MPC, the optimal control problem (3.4) incorporates three central design ingredients: the terminal constraint set $\mathcal{X}_f \subseteq \mathcal{X}$, an auxiliary local control law $\pi_f : \mathcal{X}_f \rightarrow \mathcal{U}$, and the terminal cost function $V_f : \mathcal{X}_f \rightarrow \mathbb{R}_{\geq 0}$ [2, Assumption 2.14], [3, Assumption 5.9].

Specifically, these terminal ingredients must satisfy state and input constraints, positive invariance of the terminal set, and local Lyapunov decrease of the terminal cost:

$$\mathcal{X}_f \subseteq \mathcal{X}, \quad (3.7a)$$

$$\pi_f(x) \in \mathcal{U}, \quad \forall x \in \mathcal{X}_f, \quad (3.7b)$$

$$f(x, \pi_f(x)) \in \mathcal{X}_f, \quad \forall x \in \mathcal{X}_f, \quad (3.7c)$$

$$V_f(f(x, \pi_f(x))) - V_f(x) \leq -\ell(x, \pi_f(x)), \quad \forall x \in \mathcal{X}_f. \quad (3.7d)$$

The conditions (3.7a) and (3.7b) ensure that terminal states and their corresponding local control input remain within the admissible state and input bounds $\mathcal{X}$ and $\mathcal{U}$. Additionally, positive invariance in (3.7c) ensures that every state in $\mathcal{X}_f$ remains in $\mathcal{X}_f$ under the local controller $\pi_f(x)$ [2, Def. 2.9]. This property is key for proving recursive feasibility, as it allows the terminal state of a feasible solution at horizon N to be extended by the local control input $\pi_f(x_{N|k}^*)$. This yields a feasible candidate trajectory for the subsequent sampling step $k+1$, as further described in Section 3.2.3 [2, Sec. 2.4.2]. The terminal cost decrease condition in (3.7d) acts as a local Lyapunov condition within $\mathcal{X}_f$. While positive invariance ensures that the local trajectory remains in $\mathcal{X}_f$, the decrease condition requires the terminal cost V_f to decrease along the local closed-loop trajectory by at least the stage cost $\ell(x, \pi_f(x))$ [3, Assumption 5.9].

The predictive mechanism and the geometric role of the terminal set are illustrated in Fig. 3.1. While intermediate predicted states are required to remain in the safe state space $x_{i|k} \in \mathcal{X}$ for $i = 0, \dots, N-1$, the terminal constraint enforces $x_{N|k} \in \mathcal{X}_f$, where the local controller satisfies the Lyapunov decrease condition in (3.7d) [2, Theorem 2.19].

For nonlinear systems, the terminal ingredients are typically designed by linearizing the system dynamics around the equilibrium point $(x, u) = (\mathbf{0}, \mathbf{0})$, yielding the linearized discrete-time model:

$$x_{k+1} = Ax_k + Bu_k, \quad (3.8)$$

where $A = \frac{\partial f}{\partial x}(\mathbf{0}, \mathbf{0})$ and $B = \frac{\partial f}{\partial u}(\mathbf{0}, \mathbf{0})$ represent the system Jacobians [2, Sec. 3.5]. Assuming (A, B) is stabilizable and $(A, Q^{1/2})$ is detectable, the matrix $P \succ 0$ is obtained as the unique stabilizing solution to the discrete-time Algebraic Riccati Equation (DARE) [2, Eq. (1.18)], [3, Eq. (5.23)]:

$$P = A^\top P A - A^\top P B \left(R + B^\top P B \right)^{-1} B^\top P A + Q. \quad (3.9)$$

Using P , the local linear feedback law $\pi_f(x) = -Kx$ is parameterized via the optimal Linear-Quadratic Regulator (LQR) gain [2, Eq. (1.18)], [3, Eq. (5.24)]:

$$K = \left(R + B^\top P B \right)^{-1} B^\top P A, \quad (3.10)$$

and the terminal cost function is specified quadratically as $V_f(x) = x^\top P x$. Accordingly, the terminal set $\mathcal{X}_f$ is defined in this specific case as an ellipsoidal sub-level set of this cost function [2, Sec. 3.5]:

$$\mathcal{X}_f := \left\{ x \in \mathbb{R}^{n_x} \mid x^\top P x \leq \rho \right\}. \quad (3.11)$$

By choosing the scaling parameter $\rho > 0$ sufficiently small, $\mathcal{X}_f$ satisfies both state and input constraints $\mathcal{X}$ and $\mathcal{U}$, while ensuring that the quadratic Lyapunov decrease dominates higher-order linearization errors of the underlying nonlinear dynamics.




Figure 3.1: Comparison of the terminal set $\mathcal{X}_f$ (red) with the entire state space $\mathcal{X}$ (blue).

3.2.3 Recursive Feasibility and Closed-Loop Stability

Due to state and input constraints, as well as system dynamics, the optimal control problem (OCP) may be infeasible for arbitrary states $x \in \mathcal{X}$. Therefore, the MPC domain is restricted to the *feasible set*, which is defined as [3, Sec. 3.3]:

$$\mathcal{X}_N := \left\{ x_0 \in \mathcal{X} \mid \exists (u_i)_{i=0}^{N-1} \in \mathcal{U}^N : x_{i+1} = f(x_i, u_i) \in \mathcal{X} \forall i = 0, \dots, N-1, x_N \in \mathcal{X}_f \right\}. \quad (3.12)$$

For any state $x_k \in \mathcal{X}_N$, solving (3.4) yields the optimal input sequence $(u_{i|k}^*)_{i=0}^{N-1}$ and the corresponding optimal state sequence $(x_{i|k}^*)_{i=0}^N$, where $x_{0|k}^* = x_k$ and $x_{N|k}^* \in \mathcal{X}_f$. Under the MPC policy π_0 from (3.6), the nominal closed-loop system yields the successor state $x_{k+1} = f(x_k, \pi_0(x_k)) = x_{1|k}^*$.

Recursive Feasibility with Sequence Shifting. In the nominal closed-loop setup, an MPC formulation is said to be *recursively feasible* if (3.4) is feasible for $x_k \in \mathcal{X}_N$, resulting in a subsequent state x_{k+1}

that is also feasible. This can be demonstrated with the shifted candidate input sequence $(\tilde{u}_{i|k+1})_{i=0}^{N-1}$ constructed for the state $x_{k+1} = x_{1|k}^*$, where the terminal controller π_f is appended:

$$(\tilde{u}_{i|k+1})_{i=0}^{N-1} := \left(u_{1|k}^*, \dots, u_{N-1|k}^*, \pi_f(x_{N|k}^*) \right). \quad (3.13)$$

The corresponding predicted state sequence is given by

$$(\tilde{x}_{i|k+1})_{i=0}^N := \left(x_{1|k}^*, \dots, x_{N|k}^*, f(x_{N|k}^*, \pi_f(x_{N|k}^*)) \right). \quad (3.14)$$

For $i = 0, \dots, N-2$, the pairs $(\tilde{x}_{i|k+1}, \tilde{u}_{i|k+1}) = (x_{i+1|k}^*, u_{i+1|k}^*)$ are feasible by construction. Additionally, for $i = N-1$, the input $\tilde{u}_{N-1|k+1} = \pi_f(x_{N|k}^*)$ is admissible according to (3.7b), since $x_{N|k}^* \in \mathcal{X}_f$. The computed terminal state satisfies $\tilde{x}_{N|k+1} = f(x_{N|k}^*, \pi_f(x_{N|k}^*)) \in \mathcal{X}_f \subseteq \mathcal{X}$ due to the positive invariance of $\mathcal{X}_f$ in (3.7c). Thus, the control sequence $(\tilde{u}_{i|k+1})_{i=0}^{N-1}$ is feasible for x_{k+1} , which proves that $x_{k+1} \in \mathcal{X}_N$. Consequently, for all future time steps $k \geq 0$, the closed-loop state remains in the feasible set $\mathcal{X}_N$, which confirms both positive invariance of $\mathcal{X}_N$ under π_0 and recursive feasibility.

Lyapunov Decrease of the Optimal Value Function. As the shifted candidate input sequence is feasible for x_{k+1} , the optimal cost satisfies $V_N(x_{k+1}) \leq J(x_{k+1}, (\tilde{u}_{i|k+1})_{i=0}^{N-1})$. Subtracting the optimal cost at step k from the candidate cost at time step $k+1$ yields:

$$\begin{aligned} & J(x_{k+1}, (\tilde{u}_{i|k+1})_{i=0}^{N-1}) - V_N(x_k) \\ &= \sum_{i=1}^{N-1} \ell(x_{i|k}^*, u_{i|k}^*) + \ell(x_{N|k}^*, \pi_f(x_{N|k}^*)) + V_f(f(x_{N|k}^*, \pi_f(x_{N|k}^*))) \\ &\quad - [\ell(x_k, u_{0|k}^*) + \sum_{i=1}^{N-1} \ell(x_{i|k}^*, u_{i|k}^*) + V_f(x_{N|k}^*)] \\ &= -\ell(x_k, u_{0|k}^*) + \underbrace{[V_f(f(x_{N|k}^*, \pi_f(x_{N|k}^*))) - V_f(x_{N|k}^*) + \ell(x_{N|k}^*, \pi_f(x_{N|k}^*))]}_{\leq 0 \text{ due to (3.7d)}}. \end{aligned} \quad (3.15)$$

Applying the local Lyapunov decrease condition (3.7d) ensures that the bracketed term is non-positive [3, Assumption 5.9]. Together with $V_N(x_{k+1}) \leq J(x_{k+1}, (\tilde{u}_{i|k+1})_{i=0}^{N-1})$, the optimal value function satisfies the Lyapunov decrease condition along the nominal closed-loop trajectories [3, Lemma 5.12]:

$$V_N(x_{k+1}) - V_N(x_k) \leq -\ell(x_k, \pi_0(x_k)), \quad \forall x_k \in \mathcal{X}_N. \quad (3.16)$$

Asymptotic Stability and Region of Attraction. To prove asymptotic stability of the origin, the optimal value function V_N must satisfy the criteria of a Lyapunov function on the feasible set $\mathcal{X}_N$ [3, Sec. 2.3]. Given that $Q \succ 0$, the minimum eigenvalue satisfies $\lambda_{\min}(Q) > 0$. Together with $u^\top R u \geq 0$, this results in a strict lower bound on the stage cost (3.5):

$$\ell(x, u) \geq x^\top Q x \geq \lambda_{\min}(Q) \|x\|^2 > 0, \quad \forall x \neq \mathbf{0}. \quad (3.17)$$

Consequently, this provides a negative upper bound on the Lyapunov decrease in (3.16):

$$V_N(x_{k+1}) - V_N(x_k) \leq -\ell(x_k, \pi_0(x_k)) \leq -\lambda_{\min}(Q) \|x_k\|^2 < 0, \quad \forall x_k \in \mathcal{X}_N \setminus \{\mathbf{0}\}. \quad (3.18)$$

The optimal value function satisfies the lower bound $V_N(x) \geq \ell(x, \pi_0(x)) \geq \lambda_{\min}(Q)\|x\|^2$, thus ensuring that $V_N(x) > 0$ for all $x \neq \mathbf{0}$, where $V_N(\mathbf{0}) = 0$. Under the assumption that V_N is continuous at the origin, V_N is strictly positive definite and radially bounded by class $\mathcal{K}_\infty$ functions on $\mathcal{X}_N$ [3, Def. 2.13].

Together with the positive invariance of $\mathcal{X}_N$ proven by recursive feasibility, V_N serves as a valid Lyapunov function. Therefore, the origin $\mathbf{0}$ of the nominal closed-loop system is asymptotically stable, and any trajectory initialized at $x_0 \in \mathcal{X}_N$ converges to the origin, making the feasible set $\mathcal{X}_N$ the *region of attraction* [2, Theorem 2.19], [3, Theorem 5.13].

4 Foundations of Neural Control Policies

In imitation learning, artificial neural networks (ANNs) approximate complex control laws offline to bypass the computational overhead of online Model Predictive Control (MPC) evaluations. As the machine learning foundation for this work, Section 4.1 introduces the basic architecture of Multi-Layer Perceptrons (MLPs), loss functions, backpropagation, and gradient-based optimization methods. Afterwards, Section 4.2 presents adversarial example generation and Projected Gradient Descent (PGD), which serve to identify counterexamples to the Lyapunov conditions.

4.1 Artificial Neural Networks

4.1.1 Multi-Layer Perceptrons

Multi-layer perceptrons (MLPs) are fully connected feed-forward networks that map an input through $L - 1$ hidden layers to an output layer [4, Sec. 1.4], [5, Sec. 3.1]. Let $n_l \in \mathbb{N}$ denote the number of neurons in layer $l \in \{0, \dots, L\}$, where n_0 represents the input dimension and n_L the output dimension. For an input vector $x \in \mathbb{R}^{n_0}$, forward propagation is defined recursively for layers $l = 1, \dots, L$:

$$a^{(l)} = \sigma^{(l)}(z^{(l)}) \quad \text{with} \quad z^{(l)} = W^{(l)}a^{(l-1)} + b^{(l)}, \quad \forall l = 1, \dots, L, \quad (4.1)$$

where $a^{(0)} := x$ is the network input, $z^{(l)} \in \mathbb{R}^{n_l}$ is the pre-activation vector of layer l , and $a^{(L)} := \hat{y} \in \mathbb{R}^{n_L}$ is the final network prediction. The matrix $W^{(l)} \in \mathbb{R}^{n_l \times n_{l-1}}$ and vector $b^{(l)} \in \mathbb{R}^{n_l}$ denote the weight matrix and bias vector of layer l , respectively, while $\sigma^{(l)}(\cdot)$ represents a non-linear activation function applied element-wise to the pre-activation vector $z^{(l)}$ [5, Sec. 3.2]. Without these non-linearities, the layers would simply be reduced to a single linear transformation, drastically limiting the model's ability to capture complexity [4, Sec. 1.5]. Common activation functions in neural control include the hyperbolic tangent (4.2a) and the rectified linear unit (ReLU) (4.2b), defined respectively as [5, Sec. 3.5].

$$\tanh z = \frac{e^{2z} - 1}{e^{2z} + 1} \quad (4.2a)$$

$$[z]_+ = \max(0, z) \quad (4.2b)$$

Their qualitative profiles are illustrated in Fig. 4.1.

Based on the universal approximation theorem, an MLP with at least one sufficiently sized hidden layer and non-linear activations can approximate any continuous multidimensional function on compact sets to arbitrary accuracy [5, Sec. 3.1]. The trainable parameters are represented by $\theta = \{W^{(l)}, b^{(l)}\}_{l=1}^L$. Figure 4.2 illustrates a representative architecture.




Figure 4.1: Standard activation functions for neural control policies.




Figure 4.2: Architecture of an exemplary Multi-Layer Perceptron with three inputs, two outputs and two hidden layers with four neurons each. Illustration adapted from [5, Fig. 3.1]

4.1.2 Optimization and Training

To approximate an expert control law with a neural network $\pi_\theta(x)$, the parameters θ are updated via gradient-based optimization [5, Sec. 2.5]. The objective is to find parameters θ^* that minimize the empirical training objective:

$$\theta^* \in \arg \min_{\theta} \mathcal{L}(\mathcal{D}; \theta). \quad (4.3)$$

Dataset and Primary Loss Function. The training process uses a dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^M$ containing M input-target pairs sampled across the domain. For a continuous regression tasks, a common loss term is the mean squared error (MSE) evaluated over the training dataset:

$$\mathcal{L}_{\text{MSE}}(\mathcal{D}; \theta) = \frac{1}{M} \sum_{(x,y) \in \mathcal{D}} \|\pi_\theta(x) - y\|_2^2. \quad (4.4)$$

Here, x and y denote an input sample and its corresponding target, respectively. Minimizing this objective penalizes prediction deviations, thereby encouraging the network to accurately capture the underlying mapping.

Loss Augmentation and Regularization. In supervised learning, focusing exclusively on minimizing empirical error over a finite dataset can lead to overfitting. Furthermore, it may fail to enforce underlying structural or physical properties. To improve generalization and incorporate prior knowledge directly into the optimization process, the objective is frequently augmented with weight regularization and auxiliary loss terms [5, Sec. 3.7]:

$$\mathcal{L}(\mathcal{D}; \theta) = \mathcal{L}_{\text{MSE}}(\mathcal{D}; \theta) + \omega_{\text{reg}} \mathcal{L}_{\text{reg}}(\theta) + \omega_{\text{aux}} \mathcal{L}_{\text{aux}}(\mathcal{D}; \theta). \quad (4.5)$$

where $\omega_{\text{reg}}, \omega_{\text{aux}} \geq 0$ are utilized as weights balancing prediction error, parameter regularization, and constraint satisfaction. The regularization term $\mathcal{L}_{\text{reg}}$ constrains the network's parameter space, typically realized as an ℓ_1 -norm ($\|\theta\|_1$) or ℓ_2 -norm ($\|\theta\|_2^2$) penalty. While ℓ_1 -regularization promotes sparse parameter representations, ℓ_2 -regularization penalizes large parameter magnitudes and can improve numerical stability and generalization. In practice, these penalties are applied exclusively to the network's weight matrices $W \subset \theta$, leaving bias vectors b unpenalized. [5, Sec. 3.7.2]. Alongside regularization penalties, the auxiliary loss $\mathcal{L}_{\text{aux}}$ incorporates domain knowledge into the optimization process, such as differential equations or Lyapunov conditions [5, Sec. 5], [6]. Enforcing these properties during training helps the network learn the desired physically sound solutions.

Gradient Computation via Backpropagation. For updating the network parameters θ using gradient-based optimization, the partial derivatives of the loss $\mathcal{L}(\mathcal{D}, \theta)$ with respect to all weight matrices $W^{(l)}$ and bias vectors $b^{(l)}$ must be evaluated. Backpropagation efficiently computes chain-rule derivatives of loss components that depend on the parameters through the network output by combining a forward pass (4.1) with a backward pass [5, Sec. 3.4].

During the backward pass, the algorithm computes the sensitivity of the loss with respect to the pre-activations, defined as the error signal $\delta^{(l)} = \frac{\partial \mathcal{L}}{\partial z^{(l)}}$. At the output layer L , this signal is initialized as:

$$\delta^{(L)} = \frac{\partial \mathcal{L}}{\partial a^{(L)}} \odot \frac{\partial a^{(L)}}{\partial z^{(L)}}, \quad (4.6)$$

where $\odot$ represents the element-wise (Hadamard) product. Propagating the error backward for $l = L - 1, \dots, 1$ yields the recursive relation:

$$\delta^{(l)} = \left((W^{(l+1)})^T \delta^{(l+1)} \right) \odot \frac{\partial a^{(l)}}{\partial z^{(l)}} \quad (4.7)$$

Combining these error signals with stored activations $a^{(l-1)}$ yields the network-dependent gradients of the output-dependent loss terms. Adding the explicit derivative of the parameter regularization term $\mathcal{L}_{\text{reg}}$ completes the total gradient evaluation:

$$\frac{\partial \mathcal{L}}{\partial W^{(l)}} = \delta^{(l)} \left(a^{(l-1)} \right)^T + \omega_{\text{reg}} \frac{\partial \mathcal{L}_{\text{reg}}}{\partial W^{(l)}}, \quad \frac{\partial \mathcal{L}}{\partial b^{(l)}} = \delta^{(l)}. \quad (4.8)$$

By reusing intermediate activations and error signals, backpropagation avoids redundant derivative calculations and computes parameter gradients with a computational cost similar to that of a forward pass through the network [5, Sec. 3.4], [4, Sec. 2.3].

Optimization via Mini-Batch Gradient Descent. While full-batch gradient descent evaluates the loss gradient over the entire dataset $\mathcal{D}$, this can become computationally expensive and impractical for large dataset sizes M . Conversely, single-sample stochastic gradient descent (SGD) results in high variance of the parameter updates, leading to erratic optimization trajectories. Mini-batch gradient solves this problem by dividing the dataset into smaller subsets $\mathcal{B} \subset \mathcal{D}$ of size $B = |\mathcal{B}|$. For a mini-batch, the loss $\mathcal{L}(\mathcal{B}; \theta)$ is evaluated analogously to Eq. (4.5), with $\mathcal{D}$ replaced by $\mathcal{B}$ in all data-dependent loss terms. The mini-batch gradient g typically has lower variance than the gradient obtained from a single sample, while mini-batch processing enables efficient parallelization on modern GPU hardware. During training, the parameters θ are iteratively updated along the negative gradient direction scaled by a learning rate $\alpha > 0$. This update loop is inturn executed for a total of $E \in \mathbb{N}$ epochs [4, Sec. 2.6.1], [5, Sec. 2.6].

Adaptive Moment Estimation. Although mini-batch gradient descent improves stability compared to SGD, its performance depends on a static learning rate α . In loss landscapes with noisy mini-batch gradients or challanging curvature, a fixed learning rate can cause parameter updates to oscillate in steep directions while making slow progress or none on plateaus [4, Sec. 4.5.4].

To resolve these limitations, the Adaptive Moment Estimation (Adam) optimizer [7] computes individual, adaptive learning rates for each parameter by tracking exponential moving averages of both past gradients m_t and past squared gradients v_t [4, Sec. 4.5.5.2]:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad (4.9a)$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2, \quad (4.9b)$$

where $\beta_1, \beta_2 \in [0, 1)$ are the exponential decay rates, and g_t^2 denotes the element-wise square of the current mini-batch gradient $g_t = \nabla_{\theta} \mathcal{L}(\mathcal{B}_t; \theta_{t-1})$.

Since m_0 and v_0 are initialized as zero vectors, both moments are biased towards zero. This is counteracted by dynamic correction factors $\hat{m}_t$ and $\hat{v}_t$, which are applied at each timestep t :

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}. \quad (4.10)$$

The network parameters are then updated element-wise according to:

$$\theta_t = \theta_{t-1} - \alpha \cdot \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}, \quad (4.11)$$

where $\epsilon > 0$ is a small constant used for numerical stability. Following [7], the default hyperparameters are chosen as $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, and $\epsilon = 10^{-8}$, providing robust convergence across various deep learning architectures.

4.2 Adversarial Examples and Projected Gradient Methods

Deep neural networks are known to be vulnerable to small, intentionally crafted perturbations that drastically affect model outputs or maximize loss metrics [8], [9]. This vulnerability arises from the fundamental difference between average-case and worst-case generalization: Random sampling may fail to reliably expose inputs associated with large loss values, whereas gradient-guided optimization explicitly searches for such directions [4, Sec. 12.4]. By applying gradient ascent directly to the inputs under the constraint of a bounded ℓ_p -norm $\|\delta\|_p \leq \epsilon_\delta$, the approach targets regions where the objective value is large. In neural control, such adversarial search strategies can be adapted to serve as empirical falsification tools. Rather than aiming to deceive a model, the goal is to identify input configurations that expose potential violations of safety or stability. This provides a computationally efficient approach to find high violations during training [6].

In supervised learning, the generation of adversarial examples can be formulated as a constrained optimization problem [9], [10]. Given a trained model parameterized by θ , a clean input $x \in \mathbb{R}^{n_0}$, and its ground-truth target $y \in \mathbb{R}^{n_L}$, a white-box adversary searches for an optimal perturbation vector $\delta \in \mathbb{R}^{n_0}$ that maximizes the model's loss function $\mathcal{J}(x + \delta, y; \theta)$ [10]:

$$\max_{\delta} \quad \mathcal{J}(x + \delta, y; \theta) \quad \text{s.t.} \quad \|\delta\|_p \leq \epsilon_\delta, \quad (4.12)$$

where $\epsilon_\delta > 0$ bounds the perturbation size under an ℓ_p -norm, with typical choices being $p \in \{2, \infty\}$ [9], [10]. In practical applications, the perturbed input may additionally be clipped to the valid input range.

4.2.1 Projected Gradient Descent

To solve the loss maximization problem (4.12), the Fast Gradient Sign Method (FGSM) [9] provides a fast, linear one-step approximation by taking a single step along the sign of the loss gradient. However, in highly non-convex loss landscapes, one-step attacks may yield weaker loss-maximizing candidates than iterative attacks [10].

Projected Gradient Descent (PGD) extends FGSM into an iterative multi-step gradient ascent scheme. Following the robust optimization framework proposed by Madry [10], $\mathcal{P}$ denotes the set of allowed perturbations. For an ℓ_∞ -norm constraint, $\mathcal{P}$ forms a hyper-rectangular box centered at the origin:

$$\mathcal{P} = \{\delta \in \mathbb{R}^{n_0} \mid \|\delta\|_\infty \leq \varepsilon_\delta\}. \quad (4.13)$$

Instead of initializing the search at the clean input x , the initial candidate $x_{\text{adv}}^{(0)}$ is sampled uniformly from $x + \mathcal{P}$ to prevent the optimization from getting stuck in suboptimal local maxima [10]. At each step $j \in \{0, \dots, I-1\}$, the adversarial sample is updated along the gradient direction:

$$x_{\text{adv}}^{(j+1)} \leftarrow \Pi_{x+\mathcal{P}} \left(x_{\text{adv}}^{(j)} + \alpha_{\text{PGD}} \text{sign} \left(\nabla_{x_{\text{adv}}^{(j)}} \mathcal{J} \left(x_{\text{adv}}^{(j)}; y; \theta \right) \right) \right), \quad (4.14)$$

where $\alpha_{\text{PGD}} > 0$ denotes the step size and $I \in \mathbb{N}$ represents the total number of iterations. Furthermore, $\Pi_{x+\mathcal{P}}(\cdot)$ projects the updated sample back onto the constraint set [10]. This projection can be implemented by clipping each input component to the interval $[x - \varepsilon_\delta, x + \varepsilon_\delta]$. As a result, $x_{\text{adv}}^{(j+1)}$ satisfies the perturbation constraint at every iteration.

4.2.2 Adaptation to Neural Lyapunov Control

When moving from conventional adversarial attacks to neural Lyapunov control, the adversary's objective fundamentally changes. Rather than deceiving a classifier, the goal becomes to identify counterexample states x_{adv} that violate control-theoretic guarantees, such as Lyapunov decrease (3.3) or domain invariance (3.7c) [6].

Instead of maximizing a classification loss [9], [10], the gradient-based search is utilized to identify states associated with large violations of the Lyapunov conditions. In the setting further defined in Section 6.1 and Section 6.2.1, the Lyapunov candidate and the controller are parameterized independently as $V_\phi(x)$ and $\pi_\theta(x)$, respectively. Adapting the robust optimization paradigm [10] to deterministic control, their joint training can be formulated as a saddle-point problem:

$$\min_{(\phi, \theta)} \max_{x_{\text{adv}} \in \mathcal{X}} \mathcal{J}(x_{\text{adv}}; \phi, \theta) \quad (4.15)$$

Here, the inner maximization searches for candidate states within the training domain $\mathcal{X}$ that yield large violations of the Lyapunov conditions. The objective $\mathcal{J}(x_{\text{adv}}; \phi, \theta)$ therefore depends on both the Lyapunov-function parameters ϕ and the controller parameters θ . Conversely, the parameters are updated during the outer optimization, to resolve these violations [10].

Since uniform sampling may fail to discover violations in high-dimensional spaces [4, Sec. 12.4], gradient-guided ascent along $\nabla_x \mathcal{J}$ can focus the search on candidate states associated with large Lyapunov condition violations. In this context, PGD serves as an adversarial training mechanism, iteratively feeding high-violation candidate states back into the optimization to reinforce Lyapunov stability margins [6], [10].

5 Formal Verification of Neural Control Policies

only overview because we only use it.

5.1 Incomplete Verification

5.1.1 CROWN

5.1.2 α -CROWN

5.2 Complete Verification

5.2.1 β -CROWN

6 Learning Lyapunov-Stable Imitation Controllers

This chapter presents the methodology for learning and certifying a neural approximation of the MPC controller. The MPC policy introduced in Section 3.2.1 serves as an expert, generating state-input pairs for supervised policy training. As described in Section 6.1, this allows the neural policy to reproduce the expert control actions over the considered state domain. Subsequently, a neural Lyapunov candidate is learned to enable the stability certification of the controller. In the co-training setting, the policy parameters may additionally be refined, while the Lyapunov architecture and training objective enforce the stability conditions detailed in Section 6.2. Finally, to formally prove that the learned controller satisfies these guarantees across the continuous state space, the networks are subjected to formal verification using α, β -CROWN [11], [12], as presented in Section 6.3.

6.1 Neural Policy Approximation and Imitation Learning

Data Generation uses the MPC expert policy π_0 defined in Section 3.2.3. To ensure solver convergence and obtain consistent expert trajectories during data generation, the shifted initialization is used [3, Sec. 12.5]. By shifting the previously computed optimal sequence $(u_{i|k}^*)_{i=0}^{N-1}$ and appending the local terminal controller $\pi_f(x_{N|k}^*)$, the optimization problem initializes from a feasible candidate trajectory. This prevents the solver from getting stuck in local minima between adjacent sampling instances. Closed-loop trajectories $x_{k+1} = f(x_k, \pi_0(x_k))$ are initialized from states sampled across the mpc sampling set $\mathcal{X}_{\text{MPC}} : \{x \in \mathcal{X} \mid \underline{x}_{\text{MPC}} \leq x \leq \bar{x}_{\text{MPC}}\}$. However, the training set is defined as $\mathcal{X}_\pi : \{x \in \mathcal{X} \mid \underline{x}_\pi \leq x \leq \bar{x}_\pi\}$. Storing the visited states and their corresponding optimal control actions yields the dataset

$$\mathcal{D}_\pi = \left\{ (x^{(i)}, u^{(i)}) \right\}_{i=1}^{N_{\mathcal{D}}}, \quad (6.1)$$

where $x^{(i)} \in \mathcal{X}_\pi$ and $u^{(i)} = \pi_0(x^{(i)})$.

Imitation Policy follows the structure proposed in [6], as it enforces true control inputs at the equilibrium $\pi_\theta(x^*) = u^*$ and satisfies the control input constraints $\underline{u} \leq \pi_\theta(x) \leq \bar{u}$ by construction. The neural policy π_θ is visualized in Fig. 6.1 and formulated as

$$\pi_\theta(x) = \text{clamp}(h_\theta(x) - h_\theta(x^*) + u^*, \underline{u}, \bar{u}). \quad (6.2)$$

where $h_\theta : \mathbb{R}^{n_x} \rightarrow \mathbb{R}^{n_u}$ is an unconstrained neural network, $u^* \in \mathcal{U}$ represents the desired control action at the equilibrium x^* , and $\text{clamp}(\cdot, \bar{\cdot}, \cdot)$ acts element-wise to enforce the input constraints.




Figure 6.1: The imitation learning policy.

Imitation Objective is a multi-objective loss, which consists of two parts: a state-weighted behavioral cloning loss and a dynamics-aware penalty. To account for different physical units and numerical scales across the state and control vectors, the loss terms are normalized using strictly positive diagonal scaling matrices

$$S_x = \text{diag}(x_{\text{scale}}) \in \mathbb{R}^{n_x \times n_x} \quad \text{and} \quad S_u = \text{diag}(u_{\text{scale}}) \in \mathbb{R}^{n_u \times n_u}. \quad (6.3)$$

Using S_x , the normalized distance of a state x from the equilibrium x^* is quantified with the ℓ_1 -norm as

$$d(x) = \frac{1}{n_x} \|S_x^{-1}(x - x^*)\|_1. \quad (6.4)$$

To assign appropriate priority during behavioral cloning, a state-dependent weight is defined as

$$w_x(x) = w_{\min} + (1 - w_{\min}) \left(\frac{1}{1 + d(x)} \right)^\eta, \quad (6.5)$$

where $w_{\min} \in (0, 1)$ is the baseline weight for states far from the equilibrium and $\eta \geq 0$ controls the decay rate.

State-Weighted Behavioral Cloning Loss penalizes the squared normalized Euclidean distance between the neural network policy and the expert targets. For a mini-batch $\mathcal{B}_\pi \subset \mathcal{D}_\pi$ of size $M_\pi = |\mathcal{B}_\pi|$, the behavioral cloning loss is defined as

$$\mathcal{L}_{\text{BC}} = \frac{1}{M_\pi n_u} \sum_{(x,u) \in \mathcal{B}_\pi} w_x(x) \|S_u^{-1}(\pi_\theta(x) - u)\|_2^2. \quad (6.6)$$

Here, the state-dependent weight $w_x(x)$ prioritizes approximation accuracy near the equilibrium.

Dynamics-Aware Constraint Loss applies a one-step penalty for successor states that leave the admissible state set $\mathcal{X}_\pi$. For each state $x \in \mathcal{B}_\pi$, the next state is evaluated under the neural policy as $x^+ = f(x, \pi_\theta(x))$. Using the element-wise ReLU function defined in (4.2b), the loss is defined as

$$\mathcal{L}_f = \frac{1}{M_\pi n_x} \sum_{x \in \mathcal{B}_\pi} \left(\|[x^+ - \bar{x}_\pi]_+\|_1 + \|[x_\pi - x^+]_+\|_1 \right). \quad (6.7)$$

Here, a one-step evaluation outside the admissible set results in a penalty for the current state, forcing the learned policy to stay inside the admissible set.

Total Imitation Loss is the combination of the state-weighted behavioral cloning and dynamics-aware constraint losses resulting in a scalar objective. Scaling each term by its respective weight, ω_{BC} and ω_f , yields the overall policy loss:

$$\mathcal{L}_\pi = \omega_{\text{BC}}\mathcal{L}_{\text{BC}} + \omega_f\mathcal{L}_f. \quad (6.8)$$

By minimizing $\mathcal{L}_\pi$, the policy π_θ learns to mimic the expert policy π_0 , while penalizing state constraint violations in $\mathcal{X}_\pi$. In combination with the policy architecture in (6.2), π_θ satisfies the input constraints $\underline{u} \leq \pi_\theta(x) \leq \bar{u}$ for all states $x \in \mathbb{R}^{n_x}$ and guarantees exact equilibrium control $\pi_\theta(x^*) = u^*$. Additionally, the state-dependent weighting in (6.5) prioritizes approximation accuracy near the equilibrium, which may be beneficial establishing Lyapunov stability. Now, this imitation learning pipeline establishes an accurate expert imitation, serving as an initial policy for the subsequent Lyapunov co-training.

6.2 Lyapunov Learning for Stability Certification

To establish formal stability guarantees for learned controllers, a neural Lyapunov candidate is trained alongside the policy. A falsification-guided framework is used to find stability violations and therefore make the Lyapunov candidate a valid certificate for stability.

lyapunov and imitation policy cotraining

6.2.1 Lyapunov Network Architecture

Following [6], the Lyapunov candidate $V_\phi(x)$ is parameterized to be positive definite by construction:

$$V_\phi(x) = |g_\phi(x) - g_\phi(x^*)| + \left\| (\varepsilon I + R_\phi^\top R_\phi)(x - x^*) \right\|_1, \quad (6.9)$$

where $g_\phi : \mathbb{R}^{n_x} \rightarrow \mathbb{R}$ is a neural network, $x \in \mathbb{R}^{n_x}$ is the current state, $x^* \in \mathbb{R}^{n_x}$ is the equilibrium, $\varepsilon > 0$ is a small regularization constant, and $R_\phi \in \mathbb{R}^{n_x \times n_x}$ is a learnable matrix. While the first term is only positive semi-definite and may vanish away from the equilibrium, the second term guarantees strict positive definiteness $\forall x \in \mathcal{X} \setminus \{x^*\}$, as the matrix $(\varepsilon I + R_\phi^\top R_\phi)$ is strictly positive definite by construction for any $\varepsilon > 0$. The complete architectural layout of $V_\phi(x)$ is illustrated in Fig. 6.2.




Figure 6.2: The neural Lyapunov architecture.

6.2.2 Loss Formulation

add more explanations

To ensure that the Lyapunov candidate described in (6.9) can be certified for stability, the loss function must comprise multiple components. It utilizes the previously introduced ReLU function $[\cdot]_+$ for one-sided penalties, alongside a weighted mean to aggregate the sample losses. Central to the following loss terms is the estimation of the region of attraction (ROA). Throughout the training process, a target level ρ is considered, leading to the corresponding ρ -sublevel set within the training domain $\mathcal{X}_V := \{x \in \mathcal{X}_\pi \mid \underline{x}_V \leq x \leq \bar{x}_V\}$, defined as

$$\mathcal{V}_\rho := \{x \in \mathcal{X}_V \mid V_\phi(x) \leq \rho\}. \quad (6.10)$$

To prevent this certified region from trivially collapsing towards the origin during optimization, and to ensure numerical stability during normalizations, ρ is strictly constrained by a lower bound $\rho_{\min} > 0$. Furthermore, the successor states $x_V^+ = f(x, \pi_V(x))$ needed during loss computation are evaluated on-the-fly using the system dynamics and the policy. The latter may also be trained in the co-training setup and is thus referred to as π_V hereafter.

Relative Lyapunov Decrease Loss penalizes trajectories along which the Lyapunov candidate fails to decay. For each sample, the decrease violation is defined as

$$e_V = (1 - \kappa)V_\phi(x) - V_\phi(x_V^+), \quad (6.11)$$

where $\kappa \in (0, 1)$ is a hyperparameter that controls the desired rate of decrease. To balance the optimization scale across different magnitudes of V_ϕ , we normalize this violation to obtain the relative decrease loss:

$$\mathcal{L}_{\text{dec}} = \left[\frac{e_V}{\max(|V_\phi(x)|, \varepsilon_{\text{rel}})} \right]_+, \quad (6.12)$$

where $\varepsilon_{\text{rel}} > 0$ prevents division by zero. Note that the decrease normalization can be omitted depending on the training configuration.

Relative State Invariance Loss penalizes successor states leaving the Lyapunov training domain $\mathcal{X}_V$. To evaluate the relative one-step state-constraint violation, the scaling matrix $S_{\text{range}} = \text{diag}(\bar{x}_V - \underline{x}_V)$ is utilized. The loss per sample for the next state x_V^+ is computed as

$$\mathcal{L}_{\text{inv}} = \left\| S_{\text{range}}^{-1} [x_V^+ - \bar{x}_V]_+ \right\|_1 + \left\| S_{\text{range}}^{-1} [\underline{x}_V - x_V^+]_+ \right\|_1. \quad (6.13)$$

Combined ρ -Gated Condition Loss is the sample-wise ρ -gated combination of the relative Lyapunov decrease loss from (6.12) and the relative state invariance loss from (6.13),

$$\mathcal{L}_{\text{cond}} = \mathcal{L}_{\text{dec}} + \omega_{\text{inv}} \cdot \mathcal{L}_{\text{inv}}, \quad (6.14)$$

where ω_{inv} is a hyperparameter balancing the relative importance of the two loss components. However, to apply the ρ -gate appropriately, $\mathcal{L}_{\text{cond}}$ is weighted for each sample using the soft sublevel weight

$$w_\rho = \sigma \left(s \cdot \frac{\rho - V_\theta(x)}{\max(\rho, \varepsilon_{\text{rel}})} \right), \quad (6.15)$$

where $s > 0$ denotes the gate sharpness. This smooth formulation ensures that the weights inside $\mathcal{V}_\rho$ are ≈ 1 , whereas the weights for states far outside of ρ approach zero. Evaluating these terms on a mini-batch $\mathcal{B}_\rho$ drawn from the combined state pools $\mathcal{S} \cup \mathcal{C}$ (details in Section 6.2.3), the weighted mean of the combined condition loss results in the final loss formulation:

$$\mathcal{L}_{\rho\text{-gated}} = \frac{\sum_{\mathcal{B}_\rho} w_\rho \mathcal{L}_{\text{cond}}}{\sum_{\mathcal{B}_\rho} w_\rho}. \quad (6.16)$$

The complete data flow, from evaluating the batch constraints to computing this final gated loss, is summarized in Fig. 6.3.




Figure 6.3: Combined gated condition loss $\mathcal{L}_{\rho\text{-gated}}$

ROA Loss tries to enlarge the region of attraction by uniformly drawing a batch $\mathcal{B}_{\text{ROA}}$ from $\mathcal{X}_V \setminus \xi \mathcal{X}_V$, where $\xi \in [0, 1)$ is a scaling factor defining the inner boundary of the shell, as shown in Fig. 6.4. The loss is then evaluated over these candidates using [6]

$$\mathcal{L}_{\text{ROA}} = \frac{1}{|\mathcal{B}_{\text{ROA}}|} \sum_{x \in \mathcal{B}_{\text{ROA}}} \ln \left(\left[\frac{V_\phi(x)}{\rho} - 1 \right]_+ + 1 \right). \quad (6.17)$$

Consequently, the loss function penalizes shell samples that fall outside the current ρ subsets. As the estimated region of attraction expands, fewer samples are penalized, causing the loss to decrease.




Figure 6.4: Illustration of uniformly drawn samples in the shell batch $\mathcal{B}_{\text{ROA}}$ with $\xi = 0.6$.

LIRPA-Condition Loss makes the Lyapunov function more verifier-friendly by penalizing formal violations of the Lyapunov decrease condition as defined in (6.11). These violations are identified using the LIRPA framework [12]. To reduce the computational cost during training, LIRPA is applied only to a heuristic subset of regions. Specifically, a hyper-rectangular cell is selected if at least one of its 2^{n_x} corners, or optionally its center point x_{center} , lies within the sublevel set $\mathcal{V}_\rho$. For each selected region, LIRPA then provides a formal lower bound $e_{V,\text{lower}}$ on the signed Lyapunov decrease $-e_V$. After evaluation, the LIRPA-condition violation for each region is computed as

$$\mathcal{M}_{\text{L}} = \frac{\ln([-e_{V,\text{lower}}]_+ + 1)}{\max(V_\phi(x_{\text{center}}), \varepsilon_{\text{rel}})}. \quad (6.18)$$

The final LIRPA-condition loss is then computed as the weighted mean over all evaluated regions, where the weight is defined using an exponential function of the squared Euclidean distance between the region center x_{center} and the equilibrium x^* :

$$w_{\text{L}} = \exp \left(-\alpha_{\text{L}} \cdot \|x_{\text{center}} - x^*\|_2^2 \right), \quad (6.19)$$

where $\alpha_L > 0$ is a hyperparameter controlling the decay rate of the weight. Finally, the LIRPA-condition loss over N_L regions is computed as

$$\mathcal{L}_L = \frac{\sum_{i=1}^{N_L} w_L^{(i)} \cdot \mathcal{M}_L^{(i)}}{\sum_{i=1}^{N_L} w_L^{(i)}}. \quad (6.20)$$




Figure 6.5: Lirpa regions selection with $\mathcal{V}_\rho$.

Lyapunov Scaling Loss prevents numerical collapse to zero by penalizing deviations from the initial average magnitude over a quasi-Monte Carlo batch $\mathcal{B}_V \subset \mathcal{X}_V$, sampled via a low-discrepancy Sobol sequence [13], as illustrated in Fig. 6.6. The logarithm of the mean Lyapunov value over this batch is defined as

$$\ell_V = \ln \left(\frac{1}{|\mathcal{B}_V|} \sum_{x \in \mathcal{B}_V} V_\phi(x) \right), \quad (6.21)$$

leading to the scaling loss formulation

$$\mathcal{L}_V = (\ell_{V,0} - \ell_V)^2. \quad (6.22)$$

The reference value $\ell_{V,0}$ is determined once using the initial Lyapunov candidate prior to training. To prevent the Lyapunov candidate from overfitting to a static point cloud, a new batch $\mathcal{B}_V$ is resampled from $\mathcal{X}_V$ with a specified probability at each optimization step.

Policy Scaling Loss encourages the learned policy π_V to remain close to the output of the original imitation controller π , which is based on the MPC teacher. A quasi-Monte Carlo batch $\mathcal{B}_{\text{pol}} \subset \mathcal{X}_V$ is sampled using a low-discrepancy Sobol sequence [13]. To prevent overfitting to a static set of states,




Figure 6.6: Comparisson of Sobol and uniform sampling in two dimensions.

this batch also is resampled from $\mathcal{X}_V$ with a given probability each optimization step. The normalized deviation between the learned policy and the imitation controller over this batch is formulated as

$$\mathcal{L}_{\text{pol}} = \frac{\sum_{x \in \mathcal{B}_{\text{pol}}} \|\pi_V(x) - \pi(x)\|_2^2}{\sum_{x \in \mathcal{B}_{\text{pol}}} \|\pi(x)\|_2^2} \quad (6.23)$$

R -Matrix Loss regularizes the learnable matrix R_ϕ by penalizing deviations of its Frobenius norm from its initial value $\ell_{R,0}$:

$$\mathcal{L}_R = (\ell_{R,0} - \|R\|_F)^2. \quad (6.24)$$

However, $\ell_{R,0} = \|R_\phi\|_F$ is automatically computed, based on the initial $R_{\phi,0}$ matrix.

Overall Lyapunov Loss represents the combined objective function, summing all previously defined weighted components, given by

$$\mathcal{L}_{\text{total}} = \sum_m \omega_m \mathcal{L}_m \quad (6.25)$$

where the weights $\omega_m > 0$, with $m \in \{\rho\text{-gated}, \text{ROA}, \text{L}, \text{V}, \text{pol}, \text{R}\}$ are hyperparameters that balance the relative importance of the individual loss components.

6.2.3 Falsification-Guided Training

Relying solely on uniform state sampling is not sufficient to reliably detect violations of the Lyapunov condition, especially in higher-dimensional state spaces [6]. To overcome this problem, falsification-guided training actively searches for adversarial counterexamples and incorporates them into the

training distribution as described in Section 4.2.2. This is made possible by an adversarial replay buffer that reconciles the exploration of the estimated attractor region with the targeted exploitation of known error modes.

PGD Falsification is utilized to identify counterexamples that violate the Lyapunov conditions and iteratively inject them into the counterexample pool $\mathcal{C}$. Specifically, an ℓ_∞ -Projected Gradient Descent (PGD) is employed to search for adversarial states $x_{\text{adv}} \in \mathcal{X}_V$ that maximize the condition violation defined in (6.14) and (6.15). The per-sample PGD objective is defined as

$$\mathcal{J}(x; \phi, \theta) = w_\rho(x) \mathcal{L}_{\text{cond}}(x). \quad (6.26)$$

The adversarial search is initialized using the explorative pool $\mathcal{S}$, with up to 25% of the batch supplemented by existing counterexamples to encourage local refinement. Starting from an initial candidate $x_{\text{adv}}^{(0)}$, the algorithm maximizes $\mathcal{J}$ by iteratively updating candidate states along the sign of their loss gradient:

$$x_{\text{adv}}^{(j+1)} \leftarrow \Pi_{\mathcal{X}_V} \left(x_{\text{adv}}^{(j)} + \alpha_{\text{PGD}} \text{sign} \left(\nabla_{x_{\text{adv}}^{(j)}} \mathcal{J} \left(x_{\text{adv}}^{(j)}; \phi, \theta \right) \right) \right), \quad (6.27)$$

where $\alpha_{\text{PGD}} > 0$ denotes the step size and $\Pi_{\mathcal{X}_V}(\cdot)$ projects the updated sample back onto the training domain bounds $\mathcal{X}_V$. To retain the strongest violations and mitigate the effects of overshooting, the state yielding the maximum objective value across all I iterations is preserved for each candidate. Ultimately, only states strictly within $\mathcal{V}_\rho$ that exceed the violation tolerance and lie outside the origin exclusion region are retained as counterexamples.

Adversarial Replay Buffer manages the sample distribution for the ρ -gated condition loss $\mathcal{L}_{\rho\text{-gated}}$. It maintains two distinct state pools: an explorative pool $\mathcal{S}$ and a counterexample pool $\mathcal{C}$.

As outlined in Algorithm 1, the explorative state pool $\mathcal{S}$ is repopulated at each outer epoch to balance global exploration with local refinement. First, candidate states $\mathcal{B}_\mathcal{S}$ are oversampled uniformly from the training domain $\mathcal{X}_V$. From these candidates, $N_\mathcal{S}$ states are probabilistically sampled based on a sigmoid weighting function centered around an expanded sublevel boundary $m \cdot \rho$ (with margin $m > 1$). This weighting smoothly penalizes states further outside the estimated region of attraction, significantly reducing their likelihood of being selected.

Algorithm 1: Probabilistic Explorative State Pool Resampling

Input: Lyapunov function V_ϕ , current ROA estimate ρ , margin m , sharpness s , target size $N_\mathcal{S}$

Output: Updated state pool $\mathcal{S}$ of size $N_\mathcal{S}$

- 1 Generate $4N_\mathcal{S}$ candidates: $\mathcal{B}_\mathcal{S} \sim \mathcal{U}(\mathcal{X}_V)$;
 - 2 For all $x \in \mathcal{B}_\mathcal{S}$, assign weight $w(x) \leftarrow \sigma \left(s \frac{m \cdot \rho - V_\phi(x)}{\max(\rho, 10^{-9})} \right) + \varepsilon$;
 - 3 $\mathcal{S} \leftarrow$ Sample $N_\mathcal{S}$ states from $\mathcal{B}_\mathcal{S}$ with probability $P(x) \propto w(x)$;
-

The counterexample pool $\mathcal{C}$ stores the critical violations identified during the PGD falsification step. A age-tracking mechanism limits sample lifespan, preventing overfitting outdated violations. Each injected counterexample is assigned an initial age $\tau = 0$, which is incremented after every subsequent falsification phase. Once a counterexample reaches the maximum age τ_{max} , it is removed from the pool. Consequently, the condition loss continuously penalizes recent violations while discarding resolved edge cases as the Lyapunov candidate evolves.

6.2.4 ρ -Sublevel Estimation

To evaluate the ρ -gated condition loss in (6.16), the target sublevel bound ρ must be continuously estimated throughout training as the Lyapunov candidate $V_\phi(x)$ evolves. Since the goal is to identify the largest sublevel set $\mathcal{V}_\rho := \{x \in \mathcal{X}_V \mid V_\phi(x) \leq \rho\}$ that remains within the training domain $\mathcal{X}_V$, the sublevel set reaches the boundary of the domain at its maximum admissible value. Therefore, the theoretical upper bound for ρ is given by the minimum value of $V_\phi(x)$ on the boundary:

$$\rho_{\max} = \min_{x \in \partial\mathcal{X}_V} V_\phi(x). \quad (6.28)$$

Following the boundary-evaluation principle established by [6], evaluating $V_\phi(x)$ along the boundary $\partial\mathcal{X}_V$ is sufficient to estimate $\rho_{\max}$ without searching the entire domain. However, because $\partial\mathcal{X}_V$ can be a high-dimensional surface and $V_\phi(x)$ defines a non-convex landscape, simple uniform sampling on the boundary may fail to identify low-valued boundary points. Therefore, Projected Gradient Descent (PGD) is employed to search for local minima of $V_\phi(x)$ along the boundary.

Starting from uniform boundary seeds $\mathcal{B}_\partial \sim \mathcal{U}(\partial\mathcal{X}_V)$, PGD *iteratively minimizes* $V_\phi(x)$ along the negative gradient. To maintain candidate states on the hyper-rectangular boundary surface $\partial\mathcal{X}_V$, the projection operator $\Pi_{\partial\mathcal{X}_V}$ projects the updated states onto the domain bounds while keeping the active face-defining coordinate fixed at its boundary value. For an unconstrained gradient step z originating from a boundary face along dimension j with active limit $c_j \in \{\underline{x}_j, \bar{x}_j\}$, the projection is defined element-wise for $i = 1, \dots, n_x$ as:

$$[\Pi_{\partial\mathcal{X}_V}(z)]_i = \begin{cases} c_j, & \text{if } i = j, \\ \text{clamp}(z_i, \underline{x}_i, \bar{x}_i), & \text{if } i \neq j. \end{cases} \quad (6.29)$$

To encourage the region of attraction to expand during co-training, the estimated boundary minimum is intentionally overestimated by applying a growth factor $\gamma > 1$:

$$\tilde{\rho} = \max \left(\rho_{\min}, \gamma \cdot \min_{x \in \mathcal{B}_\partial} V_\phi(x) \right). \quad (6.30)$$

By targeting a value for $\tilde{\rho}$ slightly above the currently safe maximum, the formulation deliberately leads to small violations near the boundary. These violations are subsequently captured by the adversarial replay buffer, pushing the optimization to expand the sublevel set. Finally, to reduce fluctuations caused by changes in the estimated boundary minimum during training, the gating bound ρ is updated using an exponential moving average (EMA):

$$\rho \leftarrow v\rho + (1 - v)\tilde{\rho}, \quad (6.31)$$

where $v \in [0, 1)$ denotes the EMA decay rate.

6.2.5 Overall Training Procedure

The complete falsification-guided training procedure is illustrated in Fig. 6.7. It consists of an outer loop for data generation and an inner loop for network optimization.

At the beginning of each outer epoch, boundary samples $\mathcal{B}_\partial$ are used to estimate the current sublevel

bound ρ , and the ROA shell batch $\mathcal{B}_{\text{ROA}}$ is redrawn. Based on the updated estimation of ρ , the active LiRPA regions are selected and PGD-based falsification is performed within the corresponding sublevel set. The resulting counterexamples are added to $\mathcal{C}$, after which the explorative state pool $\mathcal{S}$ is resampled. The inner loop is executed K_{steps} times per outer epoch. At each step, a mini-batch $\mathcal{B}_\rho$ is sampled from $\mathcal{S} \cup \mathcal{C}$ and evaluated together with the active LiRPA regions, the ROA shell batch, and Sobol batches $\mathcal{B}_V$ and $\mathcal{B}_{\text{pol}}$. Their corresponding loss terms are combined into $\mathcal{L}_{\text{total}}$, which is minimized by backpropagation with respect to the trainable network parameters. Repeating this procedure over N_{outer} epochs progressively reduces Lyapunov-condition violations and enlarges the estimated certified region of attraction.




Figure 6.7: Overview of the falsification-guided training framework. The architecture is divided into an outer epoch-level loop responsible for state space exploration, boundary estimation, and adversarial mining, and an inner step-level loop that optimizes the network parameters based on the generated sample pools.

After completion of the training procedure, the optimized parameters of the policy π_V and the Lyapunov

candidate V_ϕ yield a candidate control-Lyapunov pair. However, while falsification-guided adversarial training seeks to eliminate counterexamples over the trained state domain, it does not guarantee the complete absence of stability violations across the continuous state space. To transition from empirical validation to a formal stability proof, the learned candidate pair (π_V, V_ϕ) is subsequently subjected to formal neural network verification, as detailed in the following section.

6.3 Formal Verification Methodology

This section adapts the search algorithm for the maximal certifiable sublevel set $\mathcal{V}_\rho$ proposed in [6]. Specifically, for a given target level $\rho > 0$, formal verification aims to certify that all states within the candidate region $\mathcal{V}_\rho := \{x \in \mathcal{B}_C \mid V_\phi(x) \leq \rho\}$ satisfy controlled regional invariance and exponential decay:

$$V_\phi(x) \leq \rho \implies \begin{cases} x^+ \in \mathcal{B}_C, \\ V_\phi(x^+) \leq (1 - \kappa)V_\phi(x), \end{cases} \quad (6.32)$$

where $\kappa \in (0, 1)$ specifies the targeted convergence rate and $x^+ = f(x, \pi_\theta(x))$ denotes the discrete one-step state transition.

Before evaluating the stability conditions, the equilibrium x^* must be explicitly excluded from the bounded certification domain $\mathcal{B}_C := \{x \in \mathbb{R}^n \mid \underline{x}_C \leq x \leq \bar{x}_C\} \subseteq \mathcal{B}_V$. By construction, $V_\phi(x^*) = 0$ and $x^+ = x^*$ hold at the equilibrium, causing the strict positivity and strict decrease conditions to evaluate to false at this exact point. Since these conditions are only required to hold for $x \neq x^*$, and positive definiteness is structurally enforced by the network architecture, formal verification is instead performed over the punctured evaluation domain $\mathcal{B}_{\text{eval}} := \mathcal{B}_C \setminus \mathcal{B}_{\text{hole}}$, where $\mathcal{B}_{\text{hole}}$ denotes a small hyper-rectangular region centered at x^* . For the solver to handle the punctured certification domain, $\mathcal{B}_{\text{eval}}$ is split into multiple subregions, as shown in Fig. 6.8.




Figure 6.8: 2D Hole exclusion example with $\mathcal{B}_{\text{hole}} := \{x \in \mathbb{R}^2 \mid \|x\|_\infty < 0.25\}$.

To certify property (6.32) over $\mathcal{B}_{\text{eval}}$ using $\alpha\beta$ -CROWN [11], [12], the implication is rewritten into a disjunctive formulation supported by the verifier: a state must either lie outside $\mathcal{V}_\rho$ or satisfy all stability conditions inside $\mathcal{B}_{\text{eval}}$. The individual logical terms are defined as follows:

$$\Phi_{\text{sub}} \equiv V_\phi(x) > \rho, \quad (6.33a)$$

$$\Phi_{\text{pos}} \equiv V_\phi(x) > 0, \quad (6.33b)$$

$$\Phi_{\text{inv}} \equiv x_C \leq x^+ \leq \bar{x}_C, \quad (6.33c)$$

$$\Phi_{\text{dec}} \equiv (1 - \kappa)V_\phi(x) > V_\phi(x^+). \quad (6.33d)$$

Here, (6.33a) excludes states outside the target sublevel set, while (6.33b) to (6.33d) explicitly enforce positive definiteness, state constraint invariance, and exponential decrease. Combining these terms yields the overall state-wise specification:

$$\Phi(x; \rho) \equiv \Phi_{\text{sub}} \vee (\Phi_{\text{pos}} \wedge \Phi_{\text{inv}} \wedge \Phi_{\text{dec}}). \quad (6.34)$$

To determine whether a candidate level ρ is valid across the entire continuous evaluation domain, the global verification condition $\Psi(\rho)$ evaluates whether $\Phi(x; \rho)$ holds for all states:

$$\Psi(\rho) \equiv \forall x \in \mathcal{B}_{\text{eval}} : \Phi(x; \rho). \quad (6.35)$$

The $\alpha\beta$ -CROWN verifier evaluates $\Psi(\rho)$ iteratively within the scaling and bisection procedure outlined in Algorithm 2. The search operates in two distinct phases: First, a scaling phase exponentially expands or contracts ρ by a factor $s > 1$ to bracket the valid sublevel set boundary within an interval $[\rho_{\text{lo}}, \rho_{\text{up}}]$. Second, a bisection phase narrows this interval until the gap falls below a tolerance ε . Upon termination, the lower bound yields the maximum certified value $\rho^* \approx \rho_{\text{lo}}$. Compared to the proposed formulation by Yang in [6], explicit iteration limits N_{scale} and N_{bisect} are integrated into both phases. This caps the maximum number of verifier calls, effectively preventing excessive computational overhead on difficult verification problems.

Algorithm 2: Certified Value Search via Scaling and Bisection

Input: $\rho_0, \rho_{\min}$, evaluator Ψ , scale $s > 1$, tol ε , limits $N_{\text{scale}}, N_{\text{bisection}}$

Output: Max certified value ρ^*

```
1  $\rho \leftarrow \max(\rho_0, \rho_{\min});$ 
  // Phase 1: Scaling
2 if  $\Psi(\rho)$  then
3   for  $k \leftarrow 1$  to  $N_{\text{scale}}$  do
4     if  $\neg \Psi(\rho)$  then
5       break;
6     end
7      $\rho_{\text{lo}} \leftarrow \rho; \quad \rho \leftarrow \rho \cdot s;$ 
8   end
9    $\rho_{\text{up}} \leftarrow \rho;$ 
10 else
11   for  $k \leftarrow 1$  to  $N_{\text{scale}}$  do
12     if  $\Psi(\rho) \vee \rho = \rho_{\min}$  then
13       break;
14     end
15      $\rho_{\text{up}} \leftarrow \rho; \quad \rho \leftarrow \max(\rho_{\min}, \rho/s);$ 
16   end
17   if  $\neg \Psi(\rho)$  then
18     return 0 ; // No valid value  $\geq \rho_{\min}$ 
19   end
20    $\rho_{\text{lo}} \leftarrow \rho;$ 
21 end
  // Phase 2: Bisection
22 for  $k \leftarrow 1$  to  $N_{\text{bisection}}$  do
23   if  $\rho_{\text{up}} - \rho_{\text{lo}} \leq \varepsilon$  then
24     break;
25   end
26    $\rho_{\text{mid}} \leftarrow \frac{1}{2}(\rho_{\text{lo}} + \rho_{\text{up}});$ 
27   if  $\Psi(\rho_{\text{mid}})$  then
28      $\rho_{\text{lo}} \leftarrow \rho_{\text{mid}};$ 
29   else
30      $\rho_{\text{up}} \leftarrow \rho_{\text{mid}};$ 
31   end
32 end
33 return  $\rho_{\text{lo}};$ 
```

7 Experimental Evaluation and Certification Results

7.1 System Modeling and Problem Formulation

For our experimental setup, two benchmark systems are selected: the double integrator and the cartpole. These systems provide a good balance between simplicity and complexity allowing us to certify the stability of the closed-loop system.

7.1.1 Double Integrator

The double integrator is a simple second-order linear system; therefore, it serves as the ideal baseline system in the learning and certification pipeline. Discretizing the continuous-time dynamics using the explicit Euler method with a sampling time of $\Delta t = 0.1$, leads to the discrete-time state-space equation:

$$x_{k+1} = f_{\text{DI},d}(x_k, u_k) = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} x_k + \begin{bmatrix} 0 \\ \Delta t \end{bmatrix} u_k \quad (7.1)$$

where $x_k = [p_k \ v_k]^\top$ is the state vector containing the position and velocity of the system, respectively, and u_k denotes acceleration.

The *MPC setup* is formulated according to the problem in (3.4) subject to the double-integrator dynamics in (7.1). A prediction horizon of $N = 20$ steps is chosen, and the states and inputs are penalized using the weighting matrices $Q = \text{diag}(15, 1)$ and $R = 0.1$.

7.1.2 Cartpole

The inverted pendulum on a cart, commonly referred to as the cartpole system, is used as a nonlinear benchmark system in this work. There are four states in the system: the cart position p , the cart velocity v , the pendulum angle θ , and the pendulum angular velocity ω . Here, the state vector is defined as $x = [p \ v \ \theta \ \omega]^\top$, and the input u is the force F applied to the cart. This leads to the following continuous-time dynamics of the cartpole system:

$$f_{\text{CP}}(x, u) = \begin{bmatrix} \dot{p} \\ \dot{v} \\ \dot{\theta} \\ \dot{\omega} \end{bmatrix} = \begin{bmatrix} v \\ \frac{F + ml\omega^2 \sin(\theta) - mg \sin(\theta) \cos(\theta)}{M + m - m \cos^2(\theta)} \\ \omega \\ \frac{-F \cos(\theta) - ml\omega^2 \sin(\theta) \cos(\theta) + (M + m)g \sin(\theta)}{l(M + m - m \cos^2(\theta))} \end{bmatrix} \quad (7.2)$$

where $M = 1$ kg is the mass of the cart, $m = 0.1$ kg is the mass of the pendulum, $l = 0.8$ m is the length of the pendulum, and $g = 9.81$ m/s² is the acceleration due to gravity. Discretizing the continuous-time dynamics using the explicit Euler method with a sampling time of $\Delta t = 0.05$ s, leads to the discrete-time state-space equation:

$$x_{k+1} = f_{\text{CP},d}(x_k, u_k) = x_k + \Delta t f_{\text{CP}}(x_k, u_k) \quad (7.3)$$

For the nonlinear benchmark, a prediction horizon of $N = 40$ stages is chosen. The system states and control inputs are penalized within the optimization framework of (3.4) using the weighting matrices $Q = \text{diag}(10, 1, 10, 0.01)$ and $R = 0.5$, respectively, subject to the dynamics defined in (7.3).

7.2 Linear System Results: Double Integrator

Legend below




Figure 7.1: Double integrator cost descent. $\Delta V = V_N(x_{k+1}) - V_N(x_k)$

7.3 Nonlinear System Results: Cartpole




Figure 7.2: Double integrator Lyapunov landscape of the MPC optimal value function.




Figure 7.3: Double integrator Lyapunov landscape and ROA: value



8 Discussion

8.1 Qualitative Discussion of Results

8.2 Challenges with Nonlinear Dynamics

8.3 Scalability and Limitations



9 Conclusion

A Appendix

A.0.1 Imitation Learning

Note: If the MPC problem becomes infeasible at any time during one trajectory, all training pairs associated with that trajectory are discarded.

To reduce the overrepresentations of regions with high sampling rates, near-duplicate state-action pairs are removed during preprocessing. The concatenated state-action vectors are normalized and mapped onto an adaptive voxel grid. For each occupied voxel, at most one representative pair is retained.

settings for abcrown

Bibliography

- [1] J. C. Butcher, *Numerical Methods for Ordinary Differential Equations*, 3rd ed. Chichester, UK: Wiley, 2016, 1 p., ISBN: 978-1-119-12150-3 978-1-119-12151-0 978-1-119-12153-4.
- [2] J. B. Rawlings, D. Q. Mayne, and M. Diehl, *Model Predictive Control: Theory, Computation and Design*, 2nd edition. Santa Barbara: Nob Hill Publishing, LLC, 2020, 623 pp., ISBN: 978-0-9759377-5-4.
- [3] L. Grüne and J. Pannek, *Nonlinear Model Predictive Control: Theory and Algorithms* (Communications and Control Engineering), 2nd ed. 2017. Cham: Springer, 2017, 456 pp., ISBN: 978-3-319-46024-6.
- [4] C. C. Aggarwal, *Neural Networks and Deep Learning: A Textbook*, Second Edition. Cham: Springer International Publishing AG, 2023, 529 pp., ISBN: 978-3-031-29642-0 978-3-031-29641-3.
- [5] S. Kollmannsberger, D. D'Angella, M. Jokeit, and L. Herrmann, *Deep Learning in Computational Mechanics: An Introductory Course* (Studies in Computational Intelligence volume 977). Cham: Springer, 2021, 104 pp., ISBN: 978-3-030-76587-3 978-3-030-76589-7.
- [6] L. Yang, H. Dai, Z. Shi, C.-J. Hsieh, R. Tedrake, and H. Zhang, "Lyapunov-stable Neural Control for State and Output Feedback: A Novel Formulation," in *Proceedings of the 41st International Conference on Machine Learning*, Vienna, Austria: PMLR 235, 2024, pp. 56 033–56 046.
- [7] D. P. Kingma and J. Ba, "Adam: A Method for Stochastic Optimization," presented at the 3rd International Conference on Learning Representations, San Diego, CA, USA: ICLR 2015, 2015.
- [8] C. Szegedy, W. Zaremba, I. Sutskever, *et al.*, "Intriguing properties of neural networks," in *2nd International Conference on Learning Representations*, Banff, AB, Canada: OpenReview.net, 2014.
- [9] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and Harnessing Adversarial Examples," in *3rd International Conference on Learning Representations*, San Diego, CA, USA, 2015.
- [10] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, "Towards Deep Learning Models Resistant to Adversarial Attacks," in *6th International Conference on Learning Representations*, Vancouver, Canada: OpenReview.net, 2018.
- [11] S. Wang, H. Zhang, K. Xu, *et al.*, "Beta-CROWN Efficient bound propagation with per-neuron split constraints for complete and incomplete neural network verification," in *Advances in Neural Information Processing Systems*, vol. 34, Curran Associates, Inc., 2021, pp. 29 909–29 921.
- [12] K. Xu, H. Zhang, S. Wang, *et al.*, "Fast and Complete: Enabling Complete Neural Network Verification with Rapid and Massively Parallel Incomplete Verifiers," in *9th International Conference on Learning Representations*, Austria: OpenReview.net, 2021.
- [13] S. Joe and F. Y. Kuo, "Constructing Sobol Sequences with Better Two-Dimensional Projections," *SIAM Journal on Scientific Computing*, vol. 30, no. 5, pp. 2635–2654, Jan. 2008, ISSN: 1064-8275, 1095-7197.